

---

# CMPE 257: Machine Learning

## ***Team Spartans:***

1. Dataset Refine and Wrangling - Rishikesh Andhare
2. Standard Muller Loop - Jack Kalavadia
3. F1 Score and insightful dashboard - Rutvik Moradiya
4. Confusion Metrices Indexes - Pramatha Nadig

## Bike Sharing Demand Prediction

### ▾ Importing all datasets

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from matplotlib.colors import ListedColormap
import seaborn as sns
from warnings import filterwarnings
filterwarnings('ignore')

pd.options.display.max_columns = None

pd.options.display.max_rows = None
pd.options.display.float_format = '{:.6f}'.format

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import classification_report
from sklearn.metrics import accuracy_score, f1_score, precision_score, recall_score
from sklearn.metrics import confusion_matrix, roc_auc_score, roc_curve, precision_recall_curve
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn import tree
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV, StratifiedKFold

import pydotplus
from IPython.display import Image
from sklearn import metrics
from sklearn.metrics import classification_report
from sklearn.metrics import cohen_kappa_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import roc_curve
from sklearn.metrics import accuracy_score

# import various functions from statsmodels
import statsmodels
import statsmodels.api as sm

from collections import Counter

plt.rcParams['figure.figsize'] = [15,8]

datasetUrl1 = 'https://drive.google.com/file/d/10wKdg9HnqQ_o9UZatlAsTMfiIpg8NtMq/'
datasetUrl1 = 'https://drive.google.com/uc?id=' + datasetUrl1.split('/')[-1]
data1 = pd.read_csv(datasetUrl1)
data1.head(5)
```

	instant	dteday	season	yr	mnt	hr	holiday	weekday	workingday	weathersit	temp	atemp	hum	windspeed	casua
0	1	2011-01-01	1	0	1	0	0	6	0	1	0.240000	0.287900	0.810000	0.000000	
1	2	2011-01-01	1	0	1	1	0	6	0	1	0.220000	0.272700	0.800000	0.000000	

```
datasetUrl2 = 'https://drive.google.com/file/d/1yVL1fUafZ5ktpPbaAAF8zQrLlBAicgN7'
datasetUrl2 ='https://drive.google.com/uc?id=' + datasetUrl2.split('/')[1]
data2 = pd.read_csv(datasetUrl2, encoding_errors='ignore')
data2.head(5)
```

	Date	Rented Bike Count	Hour	Temperature(C)	Humidity(%)	Wind speed (m/s)	Visibility (10m)	Dew point temperature(C)
0	01/12/2017	254	0	-5.200000	37	2.200000	2000	-17.600000
1	01/12/2017	204	1	-5.500000	38	0.800000	2000	-17.600000
2	01/12/2017	173	2	-6.000000	39	1.000000	2000	-17.700000
3	01/12/2017	107	3	-6.200000	40	0.900000	2000	-17.600000
4	01/12/2017	78	4	-6.000000	36	2.300000	2000	-18.600000

```
data1.dtypes
```

```
instant      int64
dteday      object
season      int64
yr          int64
mnt         int64
hr          int64
holiday     int64
weekday     int64
workingday  int64
weathersit   int64
temp        float64
atemp       float64
hum         float64
windspeed   float64
casual      int64
registered  int64
cnt         int64
dtype: object
```

```
data2.dtypes
```

```
Date      object
Rented Bike Count  int64
Hour          int64
Temperature(C)  float64
Humidity(%)    int64
Wind speed (m/s)  float64
Visibility (10m)  int64
Dew point temperature(C)  float64
Solar Radiation (MJ/m2)  float64
Rainfall(mm)      float64
Snowfall (cm)     float64
Seasons           object
Holiday           object
Functioning Day   object
dtype: object
```

```
data2.rename(columns = {'Temperature(C)': 'temp', 'Wind speed (m/s)': 'windspeed', 'Humidity(%)': 'hum', 'Rented Bike Count': 'cnt', 'Hour': 'hour'})
```

```
data2.head()
```

	dteday	cnt	hr	temp	hum	windspeed	Visibility (10m)	Dew point	temperature(C)	solarRadiation	Rainfall(mm)	Snowfa
0	01/12/2017	254	0	-5.200000	37	2.200000	2000		-17.600000	0.000000	0.000000	

## ▼ Data Preparation and wrangling

```
data1.workingday.replace( (1, 0),('Yes', 'No'), inplace=True)
```

```
data2.holiday.replace(('Holiday', 'No Holiday'), (1, 0), inplace=True)
```

```
data1.season.replace( (1,2,3,4),('Winter', 'Autumn','Summer','Spring'), inplace=True)
```

```
data2.head()
```

	dteday	cnt	hr	temp	hum	windspeed	Visibility (10m)	Dew point	temperature(C)	solarRadiation	Rainfall(mm)	Snowfa
0	01/12/2017	254	0	-5.200000	37	2.200000	2000		-17.600000	0.000000	0.000000	
1	01/12/2017	204	1	-5.500000	38	0.800000	2000		-17.600000	0.000000	0.000000	
2	01/12/2017	173	2	-6.000000	39	1.000000	2000		-17.700000	0.000000	0.000000	
3	01/12/2017	107	3	-6.200000	40	0.900000	2000		-17.600000	0.000000	0.000000	
4	01/12/2017	78	4	-6.000000	36	2.300000	2000		-18.600000	0.000000	0.000000	

```
data1=data1.drop(['instant'],axis=1)
```

```
data1.head()
```

	dteday	season	yr	mnth	hr	holiday	weekday	workingday	weathersit	temp	atemp	hum	windspeed	casual	regist
0	2011-01-01	Winter	0	1	0	0	6	No	1	0.240000	0.287900	0.810000	0.000000	3	
1	2011-01-01	Winter	0	1	1	0	6	No	1	0.220000	0.272700	0.800000	0.000000	8	
2	2011-01-01	Winter	0	1	2	0	6	No	1	0.220000	0.272700	0.800000	0.000000	5	
3	2011-01-01	Winter	0	1	3	0	6	No	1	0.240000	0.287900	0.750000	0.000000	3	
4	2011-01-01	Winter	0	1	4	0	6	No	1	0.240000	0.287900	0.750000	0.000000	0	

```
data2.head()
```

	dteday	cnt	hr	temp	hum	windspeed	Visibility (10m)	Dew point	temperature(C)	solarRadiation	Rainfall(mm)	Snowfa
0	01/12/2017	254	0	-5.200000	37	2.200000	2000		-17.600000	0.000000	0.000000	
1	01/12/2017	204	1	-5.500000	38	0.800000	2000		-17.600000	0.000000	0.000000	
2	01/12/2017	173	2	-6.000000	39	1.000000	2000		-17.700000	0.000000	0.000000	
3	01/12/2017	107	3	-6.200000	40	0.900000	2000		-17.600000	0.000000	0.000000	
4	01/12/2017	78	4	-6.000000	36	2.300000	2000		-18.600000	0.000000	0.000000	

```
data1.dteday.head()
```

```
0    2011-01-01
1    2011-01-01
2    2011-01-01
3    2011-01-01
4    2011-01-01
Name: dteday, dtype: object
```

```
pd.to_datetime(data1.dteday.head()).dt.month,
pd.to_datetime(data1.dteday.head()).dt.day,
pd.to_datetime(data1.dteday.head()).dt.year
```

```
0    2011
1    2011
2    2011
3    2011
```

```
4    2011
Name: dteday, dtype: int64
```

```
data2.head()
```

	dteday	cnt	hr	temp	hum	windspeed	Visibility (10m)	Dew point temperature(C)	solarRadiation	Rainfall(mm)	Snowfa
0	01/12/2017	254	0	-5.200000	37	2.200000	2000	-17.600000	0.000000	0.000000	
1	01/12/2017	204	1	-5.500000	38	0.800000	2000	-17.600000	0.000000	0.000000	
2	01/12/2017	173	2	-6.000000	39	1.000000	2000	-17.700000	0.000000	0.000000	
3	01/12/2017	107	3	-6.200000	40	0.900000	2000	-17.600000	0.000000	0.000000	
4	01/12/2017	78	4	-6.000000	36	2.300000	2000	-18.600000	0.000000	0.000000	

```
df = data1.merge(data2, on = ['cnt','hr','temp','hum','windspeed','dteday','holiday','workingday','season'], how = 'outer')
```

```
df.head()
```

	dteday	season	yr	mnth	hr	holiday	weekday	workingday	weathersit	temp	atemp	hum	windspeed	casual
0	2011-01-01	Winter	0.000000	1.000000	0	0	6.000000	No	1.000000	0.240000	0.287900	0.810000	0.000000	3.000000
1	2011-01-01	Winter	0.000000	1.000000	1	0	6.000000	No	1.000000	0.220000	0.272700	0.800000	0.000000	8.000000
2	2011-01-01	Winter	0.000000	1.000000	2	0	6.000000	No	1.000000	0.220000	0.272700	0.800000	0.000000	5.000000
3	2011-01-01	Winter	0.000000	1.000000	3	0	6.000000	No	1.000000	0.240000	0.287900	0.750000	0.000000	3.000000
4	2011-01-01	Winter	0.000000	1.000000	4	0	6.000000	No	1.000000	0.240000	0.287900	0.750000	0.000000	0.000000

```
df=df.drop(['registered','atemp','yr','mnth','weekday','Visibility (10m)','Dew point temperature(C)','solarRadiation','Snowfall (cm)'],axis=1)
```

```
df.head()
```

	dteday	season	hr	holiday	workingday	temp	hum	windspeed	cnt
0	2011-01-01	Winter	0	0	No	0.240000	0.810000	0.000000	16
1	2011-01-01	Winter	1	0	No	0.220000	0.800000	0.000000	40
2	2011-01-01	Winter	2	0	No	0.220000	0.800000	0.000000	32
3	2011-01-01	Winter	3	0	No	0.240000	0.750000	0.000000	13
4	2011-01-01	Winter	4	0	No	0.240000	0.750000	0.000000	1

```
pd.to_datetime(data1.dteday.head()).dt.month,
pd.to_datetime(data1.dteday.head()).dt.day,
pd.to_datetime(data1.dteday.head()).dt.year
```

```
0    2011
1    2011
2    2011
3    2011
4    2011
Name: dteday, dtype: int64
```

```
df["month"]=pd.to_datetime(df.dteday).dt.month
```

```
df["day"]=pd.to_datetime(df.dteday).dt.day
df["year"]=pd.to_datetime(df.dteday).dt.year
```

```
df.head()
```

	dteday	season	hr	holiday	workingday	temp	hum	windspeed	cnt	month	day	year
0	2011-01-01	Winter	0	0	No	0.240000	0.810000	0.000000	16	1	1	2011
1	2011-01-01	Winter	1	0	No	0.220000	0.800000	0.000000	40	1	1	2011
2	2011-01-01	Winter	2	0	No	0.220000	0.800000	0.000000	32	1	1	2011
3	2011-01-01	Winter	3	0	No	0.240000	0.750000	0.000000	13	1	1	2011

```
df=df.drop(['dteday'],axis=1)
```

```
df.dtypes
```

```

season      object
hr          int64
holiday     int64
workingday  object
temp       float64
hum        float64
windspeed  float64
cnt         int64
month       int64
day         int64
year        int64
dtype: object
```

```
df.to_csv("dataset3.csv",index=False)
```

### ▼ Remove Insignificant variables

```
df_cat = df.select_dtypes(include=[np.object])
df_cat.head()
```

	season	workingday
0	Winter	No
1	Winter	No
2	Winter	No
3	Winter	No
4	Winter	No

```
cols = list(df.columns)
```


```
sns.countplot(x='workingday',data=df_cat, palette =['pink', 'teal'])
```

```
<Axes: xlabel='workingday', ylabel='count'>
```



```
sns.countplot(x='season',data=df_cat)
```

```
<Axes: xlabel='season', ylabel='count'>
```



```
sns.heatmap(df.corr(),annot = True,color = 'y')  
plt.show()
```




## Variables Distribution

```
df.drop('cnt', axis = 1).hist(color = "b")

# adjust the subplots
plt.tight_layout()

# display the plot
plt.show()
print('Skewness:')

df.drop('cnt', axis = 1).skew()
```



## Missing Values



```
Total = df.isnull().sum().sort_values(ascending=False)
Percent = (df.isnull().sum()*100/df.isnull().count()).sort_values(ascending=False)
missing_data = pd.concat([Total, Percent], axis = 1, keys = ['Total', 'Percentage of Missing Values'])
missing_data
```

	Total	Percentage of Missing Values
season	0	0.000000
hr	0	0.000000
holiday	0	0.000000
workingday	0	0.000000
temp	0	0.000000
hum	0	0.000000
windspeed	0	0.000000
cnt	0	0.000000
month	0	0.000000
day	0	0.000000
year	0	0.000000





## Plotting variable Based on boxplots



```
df_num = df.select_dtypes(include=[np.number])
fig, ax = plt.subplots(2, 3, figsize=(15, 8))

for variable, subplot in zip(df_num.columns, ax.flatten()):
    z = sns.boxplot(x = df_num[variable], ax=subplot,color='red', orient = "h",whis=1.5) # plot the boxplot
    z.set_xlabel(variable, fontsize = 20)
```





```
lower = df_num.temp.quantile(0.25)
upper = df_num.temp.quantile(0.75)

IQR = upper - lower

# print the IQR
print(IQR)
df_num = df_num[~((df_num.temp < (lower - 1.5 * IQR)) |(df_num.temp > (upper + 1.5 * IQR)))]
```

3.24

```
lower = df_num.hum.quantile(0.25)
```

```
upper = df_num.hum.quantile(0.75)
```

```
IQR = upper - lower
```

```
print(IQR)
```

```
df_num = df_num[~((df_num.hum < (lower - 1.5 * IQR)) |(df_num.hum > (upper + 1.5 * IQR)))]
```


0.37

```
fig, ax = plt.subplots(2, 3, figsize=(15, 8))
```


```
for variable, subplot in zip(df_num.columns, ax.flatten()):
```

```
z = sns.boxplot(x = df_num[variable], ax=subplot,color='red', orient = "h",whis=1.5 )
```


```
z.set_xlabel(variable, fontsize = 20)
```




hr




holiday




temp



hum



windspeed



cnt

```
df_target = df['cnt']
df_feature = df.drop(['cnt','holiday'], axis = 1)
df_num = df_feature.select_dtypes(include = [np.number])
df_num.head()
```

	hr	temp	hum	windspeed	month	day	year
0	0	0.240000	0.810000	0.000000	1	1	2011
1	1	0.220000	0.800000	0.000000	1	1	2011
2	2	0.220000	0.800000	0.000000	1	1	2011
3	3	0.240000	0.750000	0.000000	1	1	2011
4	4	0.240000	0.750000	0.000000	1	1	2011

```
df_cat = df_feature.select_dtypes(include = [np.object])
df_cat.columns
dummy_var = pd.get_dummies(data = df_cat, drop_first = True)
```

## ▼ Data Scaling And Train\_Test Splitting

Before employing diverse classification methods to forecast student admission outcomes, let's divide the dataset into training and testing sets.

```
target = df['holiday']
X_scaler = StandardScaler()
num_scaled = X_scaler.fit_transform(df_num)

df_num_scaled = pd.DataFrame(num_scaled, columns = df_num.columns)

X = pd.concat([df_num, dummy_var], axis = 1)
#X.drop("Sold",axis = 1,inplace = True)
X.head()
```

	hr	temp	hum	windspeed	month	day	year	season_Spring	season_Summer	season_Winter	workingday_Yes
0	0	0.240000	0.810000	0.000000	1	1	2011	0	0	1	0
1	1	0.220000	0.800000	0.000000	1	1	2011	0	0	1	0
2	2	0.220000	0.800000	0.000000	1	1	2011	0	0	1	0
3	3	0.240000	0.750000	0.000000	1	1	2011	0	0	1	0
4	4	0.240000	0.750000	0.000000	1	1	2011	0	0	1	0

```
X = sm.add_constant(X)

X_train, X_test, y_train, y_test = train_test_split(X, target, random_state = 10, test_size = 0.2)

y_train = np.round(y_train).astype(int)
y_test = np.round(y_test).astype(int)

print('X_train', X_train.shape)
print('y_train', y_train.shape)
print('X_test', X_test.shape)
print('y_test', y_test.shape)

X_train (20911, 12)
y_train (20911,)
X_test (5228, 12)
y_test (5228,)
```

## ▼ Decision Tree for Classification

```
decision_tree_classifier = DecisionTreeClassifier(criterion = 'gini', random_state = 10)
decision_tree = decision_tree_classifier.fit(X_train, y_train)
```

```
!pip install -q graphviz
```

## ▼ Finding Gini Scores

Making a full decision tree model using 'entropy'.

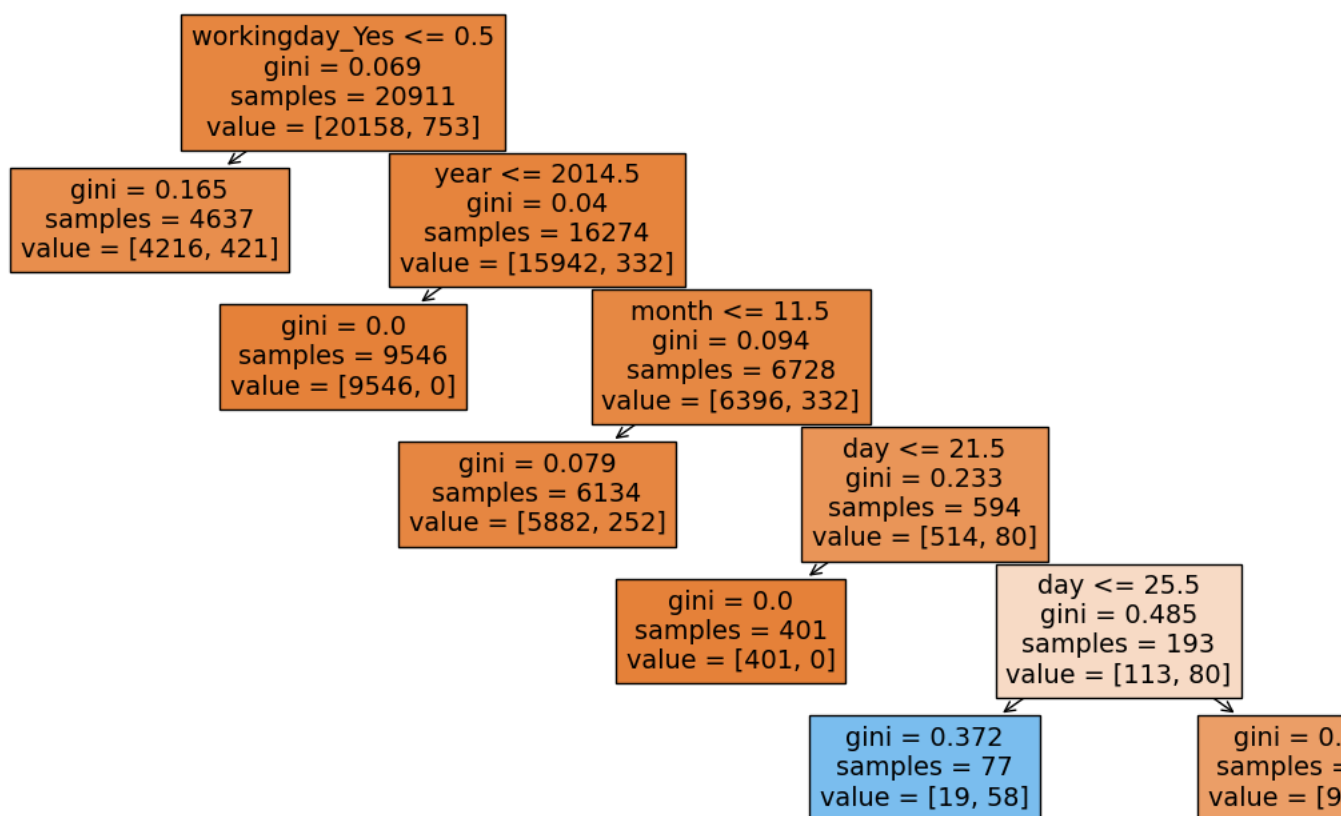
```
import pydotplus
from sklearn import tree
from IPython.display import Image
from sklearn.tree import export_graphviz

decision_tree = DecisionTreeClassifier(max_leaf_nodes=6,criterion='gini')

decision_tree = decision_tree.fit(X_train,y_train.astype(int))

labels = X_train.columns

dot_data = tree.plot_tree(decision_tree,filled = True,feature_names=X_train.columns)
```



## ▼ Train Set Performance Evaluation:

Obtaining performance metrics on the training set using a decision tree model.

```
def get_train_report(model):
    train_pred = model.predict(X_train)

    return(classification_report(y_train, train_pred))

train_report = get_train_report(decision_tree)
print(train_report)
```

	precision	recall	f1-score	support
0	0.97	1.00	0.98	20158
1	0.75	0.08	0.14	753
accuracy			0.97	20911
macro avg	0.86	0.54	0.56	20911

weighted avg	0.96	0.97	0.95	20911
--------------	------	------	------	-------

```
def get_test_report(model):  
  
    test_pred = model.predict(X_test)  
    report = classification_report(y_test, test_pred)  
    #print(report)  
    return(report, test_pred)  
  
test_report_dt, y_pred_dt= get_test_report(decision_tree)  
  
print('y_pred:', y_pred_dt)  
  
    y_pred: [0 0 0 ... 0 0 0]  
  
print(test_report_dt)
```

	precision	recall	f1-score	support
0	0.97	1.00	0.98	5049
1	0.74	0.08	0.14	179
accuracy			0.97	5228
macro avg	0.85	0.54	0.56	5228
weighted avg	0.96	0.97	0.95	5228

▼ **Hyperparameter Tuning to Mitigate Overfitting:**

**1. Hyperparameter Grid and Model Initialization:**

Defining hyperparameters for decision tree tuning to address overfitting. Initializing a decision tree model with specified hyperparameters.

**2. Model Training and Performance Evaluation:**

Training the decision tree model on the training data. Assessing model performance on both the training and test datasets.

```
tree_param_grid = {'criterion' : 'gini',  
                    'max_depth' : 5,  
                    'min_samples_split' : 4,  
                    'max_leaf_nodes' : 6,  
                    'random_state' : 10}
```

```
dt_model = DecisionTreeClassifier(criterion = 'gini',  
                                  max_depth = 5,  
                                  min_samples_split = 4,  
                                  max_leaf_nodes = 6,  
                                  random_state = 10)
```

```
decision_tree = dt_model.fit(X_train, y_train)
```

```
train_report = get_train_report(decision_tree)  
print('Train data:\n', train_report)
```


```
test_report_dt = get_test_report(decision_tree)  
print('Test data:\n', test_report_dt)
```

Train data:				
	precision	recall	f1-score	support
0	0.97	1.00	0.98	20158
1	0.75	0.08	0.14	753
accuracy			0.97	20911
macro avg	0.86	0.54	0.56	20911
weighted avg	0.96	0.97	0.95	20911

Test data:									
( '	precision	recall	f1-score	support\n\n	0	0.97	1.00	0.98	5049\n

## ▼ Features Importance

```
def FeatureImp(model):  
    imp_features = pd.DataFrame({'Features': X_train.columns,  
                                'Importance': dt_model.feature_importances_})  
  
    imp_features = imp_features.sort_values('Importance', ascending = False)  
    sns.barplot(x = 'Importance', y = 'Features', data = imp_features)  
  
    plt.title('Feature Importance', fontsize = 15)  
    plt.xlabel('Importance', fontsize = 15)  
    plt.ylabel('Features', fontsize = 15)  
  
    plt.show()  
  
FeatureImp(dt_model)
```



## ▼ 4. Random Forest for Classification Model

Initializing and training a Random Forest classifier with 10 estimators and a random state of 10.

```
rf_classification = RandomForestClassifier(n_estimators = 10, random_state = 10)  
rf_model = rf_classification.fit(X_train, y_train)
```

**Calculate performance for train set.**

```
train_report = get_train_report(rf_model)  
print(train_report)
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	20158
1	1.00	0.99	0.99	753
accuracy			1.00	20911
macro avg	1.00	0.99	1.00	20911
weighted avg	1.00	1.00	1.00	20911

#### ▼ 4.1 Tune the Hyperparameters using GridSearchCV (Random Forest)

```
tuned_paramaters = [{'criterion': ['entropy', 'gini'],
                    'n_estimators': [10, 20, 25],
                    'max_depth': [10, 15, 20],
                    'max_features': ['sqrt', 'log2'],
                    'min_samples_split': [8, 11],
                    'min_samples_leaf': [1, 5, 9],
                    'max_leaf_nodes': [2, 5, 8, 11]}]

random_forest_classification = RandomForestClassifier(random_state = 10)
rf_grid = GridSearchCV(estimator = random_forest_classification,
                      param_grid = tuned_paramaters,
                      cv = 5)
rf_grid_model = rf_grid.fit(X_train, y_train)

print('Random forest classifier top paramters list: ', rf_grid_model.best_params_,)

Random forest classifier top paramters list:  {'criterion': 'gini', 'max_depth': 10, 'max_features': 'sqrt', 'max_leaf_nodes': 5, 'min_samples_leaf': 1, 'min_samples_split': 8, 'n_estimators': 20, 'random_state': 10}

rf_model = RandomForestClassifier(criterion = rf_grid_model.best_params_.get('criterion'),
                                n_estimators = rf_grid_model.best_params_.get('n_estimators'),
                                max_depth = rf_grid_model.best_params_.get('max_depth'),
                                max_features = rf_grid_model.best_params_.get('max_features'),
                                max_leaf_nodes = rf_grid_model.best_params_.get('max_leaf_nodes'),
                                min_samples_leaf = rf_grid_model.best_params_.get('min_samples_leaf'),
                                min_samples_split = rf_grid_model.best_params_.get('min_samples_split'),
                                random_state = 10)


rf_model = rf_model.fit(X_train, y_train)
print('Classification Report for test set:\n', get_test_report(rf_model))

Classification Report for test set:
( '          precision    recall  f1-score   support\n\n         0           0.97      1.00      0.98      5049\n         1           1.00      0.99      0.99       753\n    avg / total          0.98      0.99      0.99     5802
```

#### ▼ Features Importance RF Model

```
FeatureImp(rf_model)
```

## Feature Importance



```
train_report = get_train_report(rf_model)
```

```
print(train_report)
```

	precision	recall	f1-score	support
0	0.96	1.00	0.98	20158
1	0.00	0.00	0.00	753
accuracy			0.96	20911
macro avg	0.48	0.50	0.49	20911
weighted avg	0.93	0.96	0.95	20911

season\_spring |

## ▼ Muller Classifier

... |

Setting up a comparison of various machine learning classifiers, including k-Nearest Neighbors, Support Vector Machines with linear and RBF kernels, Decision Tree, Random Forest, Neural Network, AdaBoost, and Naive Bayes.

importance

```
import numpy as np
import matplotlib.pyplot as plt
from matplotlib.colors import ListedColormap
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.datasets import make_moons, make_circles, make_classification
from sklearn.neural_network import MLPClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.gaussian_process import GaussianProcessClassifier
from sklearn.gaussian_process.kernels import RBF
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.discriminant_analysis import QuadraticDiscriminantAnalysis

names = [
    "Nearest Neighbors", "Linear SVM", "RBF SVM", "Decision Tree", "Random Forest", "#GridSearchCV",
    "Neural Net", "AdaBoost", "Naive Bayes"
]

classifiers = [
    KNeighborsClassifier(2),
    SVC(kernel="linear", C=0.025),
    SVC(gamma=2, C=1),
    DecisionTreeClassifier(max_depth=5),
    RandomForestClassifier(criterion = rf_grid_model.best_params_.get('criterion'),
                           n_estimators = rf_grid_model.best_params_.get('n_estimators'),
                           max_depth = rf_grid_model.best_params_.get('max_depth'),
                           max_features = rf_grid_model.best_params_.get('max_features'),
                           max_leaf_nodes = rf_grid_model.best_params_.get('max_leaf_nodes'),
                           min_samples_leaf = rf_grid_model.best_params_.get('min_samples_leaf'),
                           min_samples_split = rf_grid_model.best_params_.get('min_samples_split'),
                           random_state = 10),
    MLPClassifier(alpha=1, max_iter=1000),
    AdaBoostClassifier(),
    GaussianNB()]
```

```

from sklearn import metrics

def muller_classification(X_train, X_test, y_train, y_test):
    max_score = 0.0
    max_class = ''
    # iterate over classifiers
    metrics_df = pd.DataFrame({
        'Classifier': [],
        'MSE' : [],
        'MAE': [],
        'RSquared': [],
        'Test Accuracy': [],
        'Recall':[],
        'Precision': []
    })
    global m_pred
    m_pred = {}

    for name, clf in zip(names, classifiers):
        print(name)
        clf.fit(X_train, y_train)
        y_pred = clf.predict(X_test)
        score = 100.0 * clf.score(X_test, y_test)
        mean_absolute_error = np.round(metrics.mean_absolute_error(y_test, y_pred), 2)
        mean_squared_error = np.round(metrics.mean_squared_error(y_test, y_pred), 2)
        r_squared = np.round(metrics.r2_score(y_test, y_pred), 2)
        test_acc = metrics.accuracy_score(y_test, y_pred) * 100
        recall = metrics.recall_score(y_test, y_pred, average = 'weighted')
        precision = metrics.precision_score(y_test, y_pred, average = 'weighted')
        new_row = pd.DataFrame({
            'Classifier': name,
            'MSE' : mean_absolute_error,
            'MAE': mean_squared_error,
            'RSquared': r_squared,
            'Test Accuracy': test_acc,
            'Recall': recall,
            'Precision': precision}, index=[0])

        m_pred[name] = {
            'y_pred' : y_pred,
            'y_test': y_test
        }

        metrics_df = pd.concat([new_row,metrics_df.loc[:]]).reset_index(drop=True)

    print('Classifier = %s, Score (test, accuracy) = %.2f, ' %(name, score))
    if score > max_score:
        clf_best = clf
        max_score = score
        max_class = name

    print('Best Classifier = %s, Score (test, accuracy) = %.2f' %(max_class, max_score))

    return metrics_df

metrics_df = muller_classification(X_train, X_test, y_train, y_test)
metrics_df.head(10)

```



```

Nearest Neighbors
Classifier = Nearest Neighbors, Score (test, accuracy) = 98.16,
Linear SVM
Classifier = Linear SVM, Score (test, accuracy) = 96.58,
RBF SVM
Classifier = RBF SVM, Score (test, accuracy) = 96.58,
Decision Tree
Classifier = Decision Tree, Score (test, accuracy) = 96.77,
Random Forest
Classifier = Random Forest, Score (test, accuracy) = 96.58,
Neural Net
Classifier = Neural Net, Score (test, accuracy) = 96.58,
AdaBoost
Classifier = AdaBoost, Score (test, accuracy) = 96.08,

```

## ▼ Muller Loop Regressor

```

Classifier      MSE      MAE  RSquared  Test Accuracy  Recall  Precision

import numpy as np
import matplotlib.pyplot as plt
from matplotlib.colors import ListedColormap
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.neural_network import MLPRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn import metrics

names = [
    "Linear Regression",
    "MLP Regressor",
    "RandomForest Regressor",
    "Gradient Boosting Regressor",
    "KNeighbors Regressor"
]

regressors = [
    LinearRegression(),
    MLPRegressor(random_state=1, max_iter=500),
    RandomForestRegressor(max_depth=4, random_state=1),
    GradientBoostingRegressor(random_state=1),
    KNeighborsRegressor(n_neighbors=2)
]

def muller_loop(x_train, x_test, y_train, y_test):
    max_score = 0.0
    max_class = ''
    metrics_df = pd.DataFrame({
        'Regressor': [],
        'MSE' : [],
        'MAE': [],
        'RSquared': [],
        'Test Accuracy': []
    })
    for name, reg in zip(names, regressors):
        reg.fit(x_train, y_train)
        y_pred = reg.predict(x_test)
        score = 100.0 * reg.score(x_test, y_test)
        mean_absolute_error = np.round(metrics.mean_absolute_error(y_test, y_pred), 2)
        mean_squared_error = np.round(metrics.mean_squared_error(y_test, y_pred), 2)
        r_squared = np.round(metrics.r2_score(y_test, y_pred), 2)
        new_row = pd.DataFrame({
            'Regressor': name,
            'MSE' : mean_absolute_error,
            'MAE': mean_squared_error,
            'RSquared': r_squared,
            'Test Accuracy': score}, index=[0])
        metrics_df = pd.concat([new_row, metrics_df.loc[:]].reset_index(drop=True))

    print('Regressor = %s, Score (test, accuracy) = %.2f,' % (name, score))
    if score > max_score:
        reg_best = reg
        max_score = score
        max_class = name

```

```

print('***** Best Regressor = %s, Score (test, accuracy) = %.2f' %(max_class, max_score))

muller_loop(X_train, X_test, y_train, y_test )

Regressor = Linear Regression, Score (test, accuracy) = 3.75,
Regressor = MLP Regressor, Score (test, accuracy) = -82.29,
Regressor = RandomForest Regressor, Score (test, accuracy) = 14.01,
Regressor = Gradient Boosting Regressor, Score (test, accuracy) = 23.51,
Regressor = KNeighbors Regressor, Score (test, accuracy) = 63.12,
***** Best Regressor = KNeighbors Regressor, Score (test, accuracy) = 63.12

```

## Confusion matrix nd metrics

```

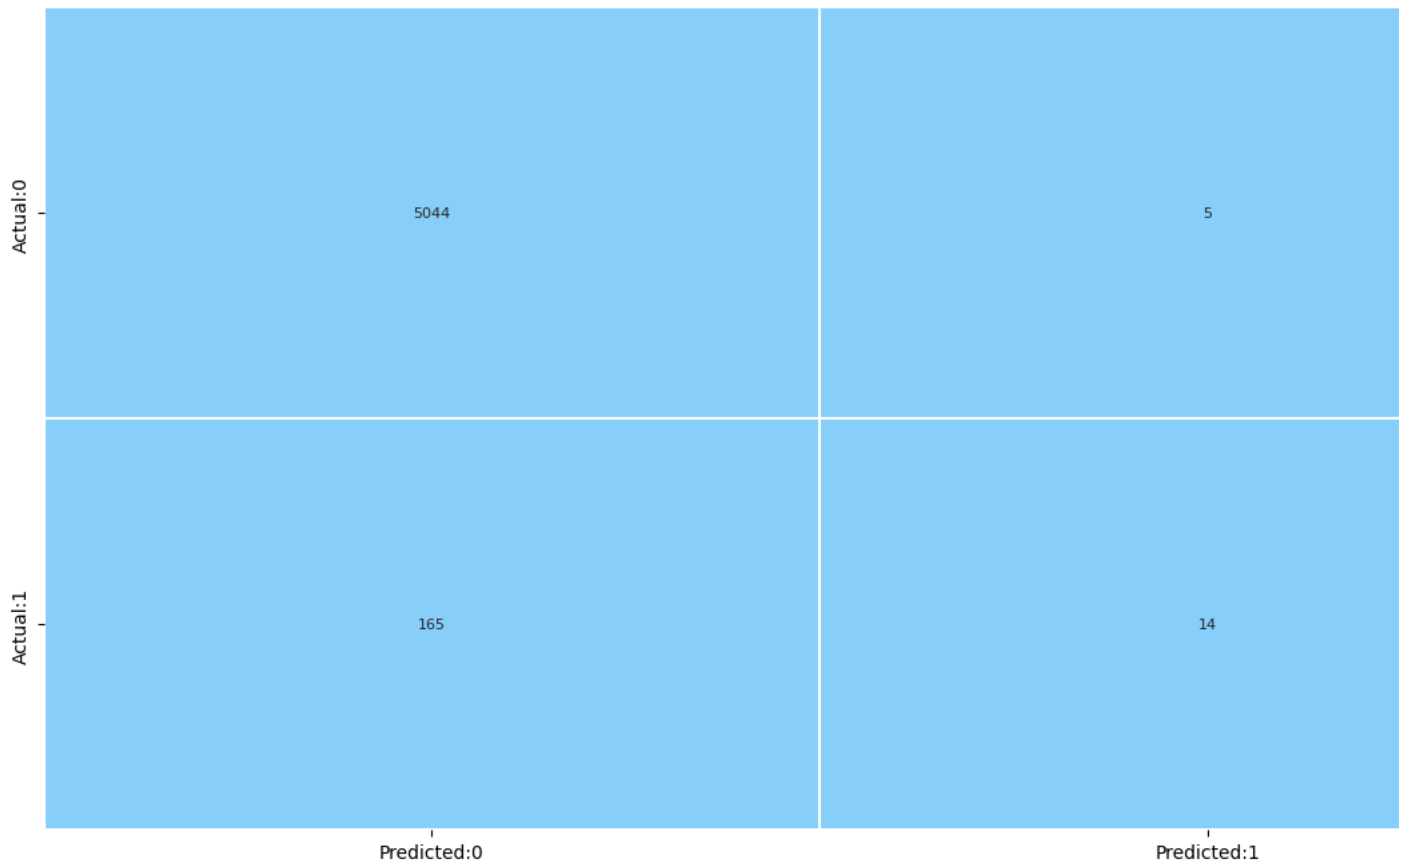
def confMatrix(y_test, y_pred):
    cm = confusion_matrix(y_test, y_pred)
    conf_matrix = pd.DataFrame(data = cm, columns = [f'Predicted:{i}' for i in range(2)], index = [f'Actual:{i}' for i in range(2)])
    return (cm, conf_matrix)

def plot_cm(conf_matrix):
    sns.heatmap(conf_matrix, annot = True, fmt = 'd', cmap = ListedColormap(['lightskyblue']), cbar = False,
                linewidths = 0.1, annot_kws = {'size':8})
    plt.xticks(fontsize = 10)
    plt.yticks(fontsize = 10)
    plt.show()

cm_dt, conf_matrix_dt = confMatrix(y_test, y_pred_dt)

plot_cm(conf_matrix_dt)

```



```

def calculateMetrics(cm, y_test, y_pred):
    # True Negatives are denoted by 'TN'
    # Actual '0' values which are classified correctly
    TN = cm[0,0]

    # True Positives are denoted by 'TP'
    # Actual '1' values which are classified correctly
    TP = cm[1,1]

```

```

# False Positives are denoted by 'FP'
# it is the type 1 error
# Actual '0' values which are classified wrongly as '1'
FP = cm[0,1]

# False Negatives are denoted by 'FN'
# it is the type 2 error
# Actual '1' values which are classified wrongly as '0'
FN = cm[1,0]

precision = TP / (TP+FP)
recall = TP / (TP+FN)
specificity = TN / (TN+FP)
f1_score = 2*((precision*recall)/(precision+recall))
accuracy = (TN+TP) / (TN+FP+FN+TP)
acc_table = classification_report(y_test, y_pred)
print(acc_table)

```

```

metrics_dt = calculateMetrics(cm_dt, y_test, y_pred_dt)

```

	precision	recall	f1-score	support
0	0.97	1.00	0.98	5049
1	0.74	0.08	0.14	179
accuracy			0.97	5228
macro avg	0.85	0.54	0.56	5228
weighted avg	0.96	0.97	0.95	5228

```

def ROC(y_test, y_pred):
    fpr, tpr, thresholds = roc_curve(y_test, y_pred)

    plt.plot(fpr, tpr)

    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.0])
    plt.plot([0, 1], [0, 1], 'r--')
    plt.title('ROC', fontsize = 15)
    plt.xlabel('False rate 1-Specificity', fontsize = 15)
    plt.ylabel('True rate Sensitivity', fontsize = 15)


    plt.text(x = 0.02, y = 0.9, s = ('AUC Score:', round(metrics.roc_auc_score(y_test, y_pred),4)))
    plt.grid(True)

```

```

ROC(y_test,y_pred_dt)

```



## ▼ Class Imbalance Handling and Evaluation

class imbalance by using SMOTE for oversampling and RandomUnderSampler for undersampling, followed by storing evaluation metrics (precision, recall, F1 score, and AUC-ROC) for various models.

```
import imblearn
from imblearn.over_sampling import RandomOverSampler, SMOTE
from imblearn.under_sampling import RandomUnderSampler, EditedNearestNeighbours
```

```
counter = Counter(y_train)
print('Before', counter)
smt = SMOTE()
```

```
X_train_sm, y_train_sm = smt.fit_resample(X_train, y_train)
```

```
counter = Counter(y_train_sm)
print('After', counter)
```

```
Before Counter({0: 20158, 1: 753})
After Counter({0: 20158, 1: 20158})
```

```
counter1 = Counter(y_train)
print('Before', counter1)
```

```
rus = RandomUnderSampler()
```

```
X_train_rus, y_train_rus = rus.fit_resample(X_train, y_train)
counter1 = Counter(y_train_rus)
print('After', counter1)
```

```
Before Counter({0: 20158, 1: 753})
After Counter({0: 753, 1: 753})
```

```
model = list()
resample = list()
precision = list()
recall = list()
F1score = list()
AUCROC = list()
```

```
def test_eval(clf_model, X_test, y_test, algo=None, sampling=None):
    y_prob=clf_model.predict_proba(X_test)
    y_pred=clf_model.predict(X_test)

    print('Confusion Matrix')
    print('='*60)
    print(confusion_matrix(y_test,y_pred),"\n")
    print('Classification Report')
    print('='*60)
    print(classification_report(y_test,y_pred),"\n")
    print('AUC-ROC')
    print('='*60)
    print(roc_auc_score(y_test, y_prob[:,1]))
```

```
model.append(algo)
precision.append(precision_score(y_test,y_pred))
recall.append(recall_score(y_test,y_pred))
F1score.append(f1_score(y_test,y_pred))
AUCROC.append(roc_auc_score(y_test, y_prob[:,1]))
resample.append(sampling)
```

## ▼ Original Data

```
estimators = [2,10,30,50,100]
max_depth = [i for i in range(5,16,2)]
min_samples_split = [2, 5, 10, 15, 20, 50, 100]
min_samples_leaf = [1, 2, 5]

tree_param_grid = {
    'max_depth':max_depth,
    'min_samples_split': min_samples_split,
    'min_samples_leaf': min_samples_leaf
}

cv = StratifiedKFold(n_splits=5, random_state=100, shuffle=True)

clf_DT = RandomizedSearchCV(dt_model, tree_param_grid, cv=cv, scoring='roc_auc', n_jobs=-1, verbose=2)
clf_DT.fit(X_train, y_train)
clf_DT.best_estimator_
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
▼
DecisionTreeClassifier
DecisionTreeClassifier(max_depth=15, max_leaf_nodes=6, min_samples_split=100,
                      random_state=10)
```

## ▼ Evaluate Decision Tree Classifier on Test Data

```
test_eval(clf_DT, X_test, y_test, 'Decision Tree', 'actual')
```

Confusion Matrix

```
=====
[[5044   5]
 [ 165  14]]
```

Classification Report

```
=====
              precision    recall  f1-score   support

     0           0.97       1.00       0.98       5049
     1           0.74       0.08       0.14        179

 accuracy              0.97       5228
 macro avg           0.85       0.54       0.56       5228
weighted avg           0.96       0.97       0.95       5228
```

AUC-ROC

```
=====
0.8117581776799654
```

```
clf_DT.fit(X_train_sm, y_train_sm)
clf_DT.best_estimator_
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
▼
DecisionTreeClassifier
DecisionTreeClassifier(max_depth=15, max_leaf_nodes=6, min_samples_split=15,
                      random_state=10)
```

```
test_eval(clf_DT, X_test, y_test, 'Decision Tree', 'smote')
```

Confusion Matrix

```
=====
[[3586 1463]
 [  20  159]]
```

Classification Report

```
=====
              precision    recall  f1-score   support

     0           0.99       0.71       0.83       5049
     1           0.10       0.89       0.18        179

 accuracy              0.72       5228
 macro avg           0.55       0.80       0.50       5228
```

weighted avg            0.96            0.72            0.81            5228

AUC-ROC

0.8315148417021567

## ▼ Random Under Sampling

```
clf_DT.fit(X_train_rus, y_train_rus)
clf_DT.best_estimator_
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
▼
DecisionTreeClassifier
DecisionTreeClassifier(max_depth=5, max_leaf_nodes=6, min_samples_leaf=5,
min_samples_split=15, random_state=10)
```

```
test_eval(clf_DT, X_test, y_test, 'Decision Tree', 'RandomUnderSampler')
```

Confusion Matrix

```
=====
[[2728 2321]
 [   1  178]]
```

Classification Report

```
=====
              precision    recall  f1-score   support

      0               1.00      0.54      0.70      5049
      1               0.07      0.99      0.13       179

 accuracy               0.56               5228
 macro avg              0.54              0.77      0.42      5228
weighted avg              0.97              0.56      0.68      5228
```

AUC-ROC

0.7813356480789934

## Random Forest

```
rf_model = RandomForestClassifier()
```

```
rf_params={'n_estimators':estimators,
           'max_depth':max_depth,
           'min_samples_split':min_samples_split}
```

```
clf_RF = RandomizedSearchCV(rf_model, rf_params, cv=cv, scoring='roc_auc', n_jobs=-1, n_iter=20, verbose=2)
clf_RF.fit(X_train, y_train)
clf_RF.best_estimator_
test_eval(clf_RF, X_test, y_test, 'Random Forest', 'actual')
```

Fitting 5 folds for each of 20 candidates, totalling 100 fits

```
Confusion Matrix
=====
[[5049    0]
 [  56  123]]
```

Classification Report

```
=====
              precision    recall  f1-score   support

      0               0.99      1.00      0.99      5049
      1               1.00      0.69      0.81       179

 accuracy               0.99               5228
 macro avg              0.99              0.84      0.90      5228
weighted avg              0.99              0.99      0.99      5228
```

AUC-ROC

0.999361563935997

▼ SMOTE Resampling and Evaluation

Random Forest Model

```
clf_RF.fit(X_train_sm, y_train_sm)
clf_RF.best_estimator_

Fitting 5 folds for each of 20 candidates, totalling 100 fits
▼ RandomForestClassifier
RandomForestClassifier(max_depth=15, min_samples_split=5, n_estimators=50)

test_eval(clf_RF, X_test, y_test, 'Random Forest', 'smote')

Confusion Matrix
=====
[[4987  62]
 [ 11 168]]

Classification Report
=====
              precision    recall  f1-score   support

    0               1.00      0.99      0.99       5049
    1               0.73      0.94      0.82        179

 accuracy               0.99       5228
 macro avg              0.86      0.96      0.91       5228
weighted avg              0.99      0.99      0.99       5228


AUC-ROC
=====
0.9966495937577107
```

Random Under Sampling

```
clf_RF.fit(X_train_rus, y_train_rus)
clf_RF.best_estimator_

Fitting 5 folds for each of 20 candidates, totalling 100 fits
▼ RandomForestClassifier
RandomForestClassifier(max_depth=13)

test_eval(clf_RF, X_test, y_test, 'Decision Tree', 'RandomUnderSampler')

Confusion Matrix
=====
[[4641  408]
 [  1 178]]

Classification Report
=====
              precision    recall  f1-score   support

    0               1.00      0.92      0.96       5049
    1               0.30      0.99      0.47        179

 accuracy               0.92       5228
 macro avg              0.65      0.96      0.71       5228
weighted avg              0.98      0.92      0.94       5228


AUC-ROC
=====
0.9922225873589658
```

Model Comparison

```
clf_compare = pd.DataFrame({'model':model,
                             'resample':resample,
                             'precision':precision,
                             'recall':recall,
```

```
        'f1-score':F1score,
        'AUC-ROC':AUCROC})

clf_compare
```

	model	resample	precision	recall	f1-score	AUC-ROC
0	Decision Tree	actual	0.736842	0.078212	0.141414	0.811758
1	Decision Tree	smote	0.098027	0.888268	0.176569	0.831515
2	Decision Tree	RandomUnderSampler	0.071228	0.994413	0.132935	0.781336
3	Random Forest	actual	1.000000	0.687151	0.814570	0.999362
4	Random Forest	smote	0.730435	0.938547	0.821516	0.996650
5	Decision Tree	RandomUnderSampler	0.303754	0.994413	0.465359	0.992223

HoloViews Data distribution

```
!pip install -q hvplot
3.2/3.2 MB 28.0 MB/s eta 0:00:00
```

```
import holoviews as hv
hv.extension('bokeh')
import hvplot.pandas
import panel as pn
from sklearn.metrics import confusion_matrix, roc_curve
```



▼ HoloViews Data Distribution Visualization with Interactive Controls:

Utilization HoloViews and Panel to create an interactive data distribution visualization with widgets for selecting **season, working day, and temperature**.

```
hv.extension('bokeh')
season_text = pn.widgets.StaticText(name='season', value='')

select_season = pn.widgets.RadioBoxGroup(
    name='season', options=['Winter', 'Summer', 'Autumn', 'Spring'],
    inline=True)

workingday_text = pn.widgets.StaticText(name='workingday', value='')
select_workingday = pn.widgets.RadioBoxGroup(
    name='workingday', options=['Yes', 'No'],
    inline=True)

select_temp = pn.widgets.IntSlider(name='temp', start= int(min(df['temp'])), end= int(max(df['temp'])), step = 1)
@pn.depends(select_season, select_workingday, select_temp)

def exp_plot(select_season, select_workingday, select_temp):
    return df[(df.season==select_season) & (df.workingday==select_workingday) & (df.temp <= select_temp)].sort_values(by='cnt').hvplot()

pn.Column(season_text, select_season, workingday_text, select_workingday, select_temp, exp_plot).embed()
```





season:

☒ Winter ☐ Summer ☐ Autumn ☐ Spring

workingday:

☐ Yes ☒ No

temp: 11



na 4 1

```
def hv_confusion_matrix(y_pred,
                        y_test,
                        title='Confusion matrix'):
    tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()
    pdf = pd.DataFrame(zip(y_pred, y_train), columns=['Prediction', 'Actual'])

    graph = pdf.groupby(['Prediction', 'Actual']).size().to_frame().reset_index()
    confusion = graph.rename(columns={0: 'Count'})
    conf_values = map(lambda l: [str(l[0]), str(l[1]), l[2]], [a.tolist() for a in confusion.values])
    return hv.HeatMap(conf_values, label='Confusion Matrix', kdims=['Predicted', 'Actual'], vdims=['Count']).sort().options(
        xrotation=45, width=500, height=500, cmap='viridis', tools=['hover'], invert_yaxis=True, zlim=(0,1))

def hv_roc_curve(y_pred, y_test):
    fpr, tpr, _ = roc_curve(y_test, y_pred)
    roc_df = pd.DataFrame(zip(fpr, tpr), columns=['Specificity', 'Sensitivity'])
    return roc_df.hvplot(x='Specificity', y='Sensitivity')

def upsample(df, n):
    df = resample(df, n_samples=n, replace=True)
    return df

def downsample(df, n):
    df = resample(df, n_samples=n, replace=False)
    return df

from sklearn.utils import resample

random = pd.concat([X, df['holiday']], axis = 1)
random = resample(random)
random_tar = random['holiday'].fillna(0)
random.drop(['holiday'], inplace=True, axis=1)
X_train, X_test, y_train, y_test = train_test_split(random, random_tar, random_state = 10, test_size = 0.2)
muller_classification(X_train, X_test, y_train, y_test)

Linear Regression
Classifier = Linear Regression, Score (test, accuracy) = 98.91,
MLP Regressor
Classifier = MLP Regressor, Score (test, accuracy) = 96.60,
RandomForest Regressor
Classifier = RandomForest Regressor, Score (test, accuracy) = 98.57,
Gradient Boosting Regressor
Classifier = Gradient Boosting Regressor, Score (test, accuracy) = 96.58,
KNeighbors Regressor
Classifier = KNeighbors Regressor, Score (test, accuracy) = 96.60,
Best Classifier = Linear Regression, Score (test, accuracy) = 98.91



|   | Classifier                  | MSE      | MAE      | RSquared  | Test Accuracy | Recall   | Precision |
|---|-----------------------------|----------|----------|-----------|---------------|----------|-----------|
| 0 | KNeighbors Regressor        | 0.030000 | 0.030000 | -0.040000 | 96.595256     | 0.965953 | 0.933064  |
| 1 | Gradient Boosting Regressor | 0.030000 | 0.030000 | -0.040000 | 96.576129     | 0.965761 | 0.951919  |
| 2 | RandomForest Regressor      | 0.010000 | 0.010000 | 0.560000  | 98.565417     | 0.985654 | 0.985864  |
| 3 | MLP Regressor               | 0.030000 | 0.030000 | -0.040000 | 96.595256     | 0.965953 | 0.933064  |
| 4 | Linear Regression           | 0.010000 | 0.010000 | 0.670000  | 98.909717     | 0.989097 | 0.988846  |



from sklearn.utils import resample

def exp_plot(event):
    temp_df = pd.concat([X, df['holiday']], axis = 1)
    temp_df = resample(temp_df)
    temp_target = temp_df['holiday'].fillna(0)
    temp_train_df = temp_df.drop(['holiday'], axis=1)
    X_train_temp, X_test_temp, y_train_temp, y_test_temp = train_test_split(temp_train_df, temp_target, random_state = 10, test_size=
```

```


hv.extension('bokeh')

samples_text = pn.widgets.StaticText(name='Select Number of Samples', value='')
select_sample_size = pn.widgets.IntSlider(name='SampleSize', start= 100, end= 506, step = 10)

button = pn.widgets.Button(name='Run Muller Loop', button_type='primary')


button.on_click(exp_plot)
pn.Column(samples_text, select_sample_size, button)

```



**Select Number of Samples:**

SampleSize: **100**



Run Muller Loop



```

hv.extension('bokeh')
model_text = pn.widgets.StaticText(name='Model', value='', styles=dict(background='red'))
select_model = pn.widgets.RadioBoxGroup( name='model', options=["Linear Regression", "MLP Regressor", "RandomForest Regressor", "(

@pn.depends(select_model)
def plot_conf_roc(select_model):
    y_test = m_pred[select_model]['y_test']
    y_pred = m_pred[select_model]['y_pred']

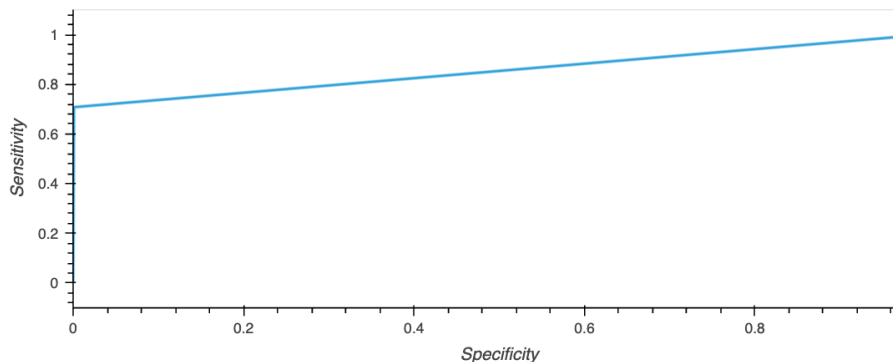
    return hv_roc_curve(y_pred,y_test) + hv_confusion_matrix(y_pred,y_test)
pn.Column(model_text, select_model, plot_conf_roc).embed()

```

**Model:**

☒ Linear Regression ☐ MLP Regressor ☐ RandomForest Regressor ☐ Gradient Boosting Regressor ☐ KNeighb



confusion matrices for bike rental prediction models, specifically for binary classification tasks with class labels "0" and "1." Users can select different regression models and evaluate their performance using these confusion matrices along with other visualizations

